

Rishabh Patil Sapid:60009200056 batch:K3

Q1)Perform RandomForest from scratch on dataset 1.

```
from random import seed
from random import randrange
from csv import reader
from math import sqrt

# Load a CSV file
def load_csv(filename):
    dataset = list()
    with open(filename, 'r') as file:
        csv_reader = reader(file)
        for row in csv_reader:
            if not row:
                continue
            dataset.append(row)
    return dataset

# Convert string column to float
def str_column_to_float(dataset, column):
    for row in dataset:
        row[column] = float(row[column].strip())

# Convert string column to integer
def str_column_to_int(dataset, column):
    class_values = [row[column] for row in dataset]
    unique = set(class_values)
    lookup = dict()
    for i, value in enumerate(unique):
        lookup[value] = i
    for row in dataset:
        row[column] = lookup[row[column]]
    return lookup

# Split a dataset into k folds
def cross_validation_split(dataset, n_folds):
    dataset_split = list()
    dataset_copy = list(dataset)
    fold_size = int(len(dataset) / n_folds)
    for i in range(n_folds):
        fold = list()
        while len(fold) < fold_size:
            index = randrange(len(dataset_copy))
            fold.append(dataset_copy.pop(index))
        dataset_split.append(fold)
    return dataset_split
```

```

# Calculate accuracy percentage
def accuracy_metric(actual, predicted):
    correct = 0
    for i in range(len(actual)):
        if actual[i] == predicted[i]:
            correct += 1
    return correct / float(len(actual)) * 100.0

# Evaluate an algorithm using a cross validation split
def evaluate_algorithm(dataset, algorithm, n_folds, *args):
    folds = cross_validation_split(dataset, n_folds)
    scores = list()
    for fold in folds:
        train_set = list(folds)
        train_set.remove(fold)
        train_set = sum(train_set, [])
        test_set = list()
        for row in fold:
            row_copy = list(row)
            test_set.append(row_copy)
            row_copy[-1] = None
        predicted = algorithm(train_set, test_set, *args)
        actual = [row[-1] for row in fold]
        accuracy = accuracy_metric(actual, predicted)
        scores.append(accuracy)
    return scores

# Split a dataset based on an attribute and an attribute value
def test_split(index, value, dataset):
    left, right = list(), list()
    for row in dataset:
        if row[index] < value:
            left.append(row)
        else:
            right.append(row)
    return left, right

# Calculate the Gini index for a split dataset
def gini_index(groups, classes):
    # count all samples at split point
    n_instances = float(sum([len(group) for group in groups]))
    # sum weighted Gini index for each group
    gini = 0.0
    for group in groups:
        size = float(len(group))
        # avoid divide by zero
        if size == 0:
            continue
        score = 0.0
        # score the group based on the score for each class

```

```

        for class_val in classes:
            p = [row[-1] for row in group].count(class_val) / size
            score += p * p
        # weight the group score by its relative size
        gini += (1.0 - score) * (size / n_instances)
    return gini

# Select the best split point for a dataset
def get_split(dataset, n_features):
    class_values = list(set(row[-1] for row in dataset))
    b_index, b_value, b_score, b_groups = 999, 999, 999, None
    features = list()
    while len(features) < n_features:
        index = randrange(len(dataset[0])-1)
        if index not in features:
            features.append(index)
    for index in features:
        for row in dataset:
            groups = test_split(index, row[index], dataset)
            gini = gini_index(groups, class_values)
            if gini < b_score:
                b_index, b_value, b_score, b_groups = index,
row[index], gini, groups
    return {'index':b_index, 'value':b_value, 'groups':b_groups}

# Create a terminal node value
def to_terminal(group):
    outcomes = [row[-1] for row in group]
    return max(set(outcomes), key=outcomes.count)

# Create child splits for a node or make terminal
def split(node, max_depth, min_size, n_features, depth):
    left, right = node['groups']
    del(node['groups'])
    # check for a no split
    if not left or not right:
        node['left'] = node['right'] = to_terminal(left + right)
        return
    # check for max depth
    if depth >= max_depth:
        node['left'], node['right'] = to_terminal(left),
to_terminal(right)
        return
    # process left child
    if len(left) <= min_size:
        node['left'] = to_terminal(left)
    else:
        node['left'] = get_split(left, n_features)
        split(node['left'], max_depth, min_size, n_features,
depth+1)
    # process right child

```

```

    if len(right) <= min_size:
        node['right'] = to_terminal(right)
    else:
        node['right'] = get_split(right, n_features)
        split(node['right'], max_depth, min_size, n_features,
depth+1)

# Build a decision tree
def build_tree(train, max_depth, min_size, n_features):
    root = get_split(train, n_features)
    split(root, max_depth, min_size, n_features, 1)
    return root

# Make a prediction with a decision tree
def predict(node, row):
    if row[node['index']] < node['value']:
        if isinstance(node['left'], dict):
            return predict(node['left'], row)
        else:
            return node['left']
    else:
        if isinstance(node['right'], dict):
            return predict(node['right'], row)
        else:
            return node['right']

# Create a random subsample from the dataset with replacement
def subsample(dataset, ratio):
    sample = list()
    n_sample = round(len(dataset) * ratio)
    while len(sample) < n_sample:
        index = randrange(len(dataset))
        sample.append(dataset[index])
    return sample

# Make a prediction with a list of bagged trees
def bagging_predict(trees, row):
    predictions = [predict(tree, row) for tree in trees]
    return max(set(predictions), key=predictions.count)

# Random Forest Algorithm
def random_forest(train, test, max_depth, min_size, sample_size,
n_trees, n_features):
    trees = list()
    for i in range(n_trees):
        sample = subsample(train, sample_size)
        tree = build_tree(sample, max_depth, min_size, n_features)
        trees.append(tree)
    predictions = [bagging_predict(trees, row) for row in test]
    return(predictions)

```

```

# Test the random forest algorithm
seed(2)
# load and prepare data
filename = 'sonar.all-data.csv'
dataset = load_csv(filename)
# convert string attributes to integers
for i in range(0, len(dataset[0])-1):
    str_column_to_float(dataset, i)
# convert class column to integers
str_column_to_int(dataset, len(dataset[0])-1)
# evaluate algorithm
n_folds = 5
max_depth = 10
min_size = 1
sample_size = 1.0
n_features = int(sqrt(len(dataset[0])-1))
for n_trees in [1, 5, 10]:
    scores = evaluate_algorithm(dataset, random_forest, n_folds,
max_depth, min_size, sample_size, n_trees, n_features)
    print('Trees: %d' % n_trees)
    print('Scores: %s' % scores)
    print('Mean Accuracy: %.3f%%' % (sum(scores)/float(len(scores))))

```

Trees: 1

Scores: [56.09756097560976, 63.41463414634146, 60.97560975609756, 58.536585365853654, 73.17073170731707]

Mean Accuracy: 62.439%

Trees: 5

Scores: [70.73170731707317, 58.536585365853654, 85.36585365853658, 75.60975609756098, 63.41463414634146]

Mean Accuracy: 70.732%

Trees: 10

Scores: [75.60975609756098, 80.48780487804879, 92.6829268292683, 73.17073170731707, 70.73170731707317]

Mean Accuracy: 78.537%

Rishabh Patil Sapid:60009200056 batch:K3

Q2) Compare the results of decision tree and random forest classifier for dataset 2 and 3

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
df=pd.read_csv('/content/Iris (2).csv')
```

```
df.head()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
df.drop(['Id'],axis=1,inplace=True)
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df["Species"] = le.fit_transform(df["Species"])
```

```
df.head()
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0

```
#decision tree
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
y=df.iloc[:,-1]
x=df.drop(['Species'],axis=1)
```

```
feature_col=list(x.columns)
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(x,y,test_size=0.2,random_state=55)
```

```

model=DecisionTreeClassifier()
model.fit(X_train,Y_train)
train_pred=model.predict(X_train)
test_pred=model.predict(X_test)
print(train_pred)

[0 0 2 2 0 1 1 0 2 1 1 0 1 0 1 2 2 2 0 2 2 0 1 2 1 0 0 1 0 0 1 2 0 2 1
 2 0
 2 0 0 0 0 2 1 0 0 2 2 2 1 1 2 2 0 0 2 1 0 1 1 0 2 2 1 1 0 1 1 2 2 0 2
 1 0
 0 0 0 2 1 2 2 2 1 0 2 1 1 2 0 2 1 2 1 1 0 2 1 0 1 0 1 0 1 2 2 1 1 0 2
 1 1
 0 2 1 1 0 1 0 2 0]

from sklearn.metrics import accuracy_score
print(accuracy_score(Y_train,train_pred))
print(accuracy_score(Y_test,test_pred))

```

```

1.0
0.9666666666666667

```

#Random Forest

```

from sklearn.ensemble import RandomForestClassifier
clf=RandomForestClassifier(n_estimators=50)
models=clf.fit(X_train,Y_train)
train_random_pred=models.predict(X_train)
test_random_pred=models.predict(X_test)
print(accuracy_score(Y_train,train_random_pred))
print(accuracy_score(Y_test,test_random_pred))

```

```

1.0
0.9666666666666667

```

#DT+FI

```

feature_col=['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm',
'PetalWidthCm']
feature_imp =
pd.Series(clf.feature_importances_,index=feature_col).sort_values(ascending=False)
feature_imp

```

```

PetalLengthCm    0.430991
PetalWidthCm     0.426522
SepalLengthCm    0.120873
SepalWidthCm     0.021614
dtype: float64

```

```
df.drop(['SepalWidthCm'],axis=1,inplace=True)
```

```
df.head()
```

	SepalLengthCm	PetalLengthCm	PetalWidthCm	Species
0	5.1	1.4	0.2	0
1	4.9	1.4	0.2	0
2	4.7	1.3	0.2	0
3	4.6	1.5	0.2	0
4	5.0	1.4	0.2	0

```
y_fi=df.iloc[:,-1]
x_fi=df.drop(['Species'],axis=1)
```

```
from sklearn.model_selection import train_test_split
X_train_fi,X_test_fi,Y_train_fi,Y_test_fi=train_test_split(x_fi,y_fi,t
est_size=0.2,random_state=55)
```

```
model=DecisionTreeClassifier()
model.fit(X_train_fi,Y_train_fi)
train_pred_fi=model.predict(X_train_fi)
test_pred_fi=model.predict(X_test_fi)
print(train_pred_fi)
```

```
[0 0 2 2 0 1 1 0 2 1 1 0 1 0 1 2 2 2 0 2 2 0 1 2 1 0 0 1 0 0 1 2 0 2 1
2 0
 2 0 0 0 0 2 1 0 0 2 2 2 1 1 2 2 0 0 2 1 0 1 1 0 2 2 1 1 0 1 1 2 2 0 2
1 0
 0 0 0 2 1 2 2 2 1 0 2 1 1 2 0 2 1 2 1 1 0 2 1 0 1 0 1 0 1 2 2 1 1 0 2
1 1
 0 2 1 1 0 1 0 2 0]
```

```
from sklearn.metrics import accuracy_score
print(accuracy_score(Y_train_fi,train_pred_fi))
print(accuracy_score(Y_test_fi,test_pred_fi))
```

```
1.0
0.9666666666666667
```

```
from sklearn.ensemble import RandomForestClassifier
clf=RandomForestClassifier(n_estimators=50)
models=clf.fit(X_train_fi,Y_train_fi)
train_random_pred_fi=models.predict(X_train_fi)
test_random_pred_fi=models.predict(X_test_fi)
print(accuracy_score(Y_train_fi,train_random_pred_fi))
print(accuracy_score(Y_test_fi,test_random_pred_fi))
```

```
1.0
0.9666666666666667
```



```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
l=[]
for i in range(1,62):
    l.append(i)
df=pd.read_csv('/content/sonar.all-data.csv',names=l)
df.head()
```

	1	2	3	4	5	6	7	8	
9 \									
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	
0.3109									
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	
0.3337									
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	
0.5598									
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	
0.0598									
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	
0.3564									
	10	...	52	53	54	55	56	57	58
\									
0	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180	0.0084
1	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140	0.0049
2	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316	0.0164
3	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050	0.0044
4	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072	0.0048
	59	60	61						
0	0.0090	0.0032	R						
1	0.0052	0.0044	R						
2	0.0095	0.0078	R						
3	0.0040	0.0117	R						
4	0.0107	0.0094	R						

[5 rows x 61 columns]

```
df.isnull().sum()
```

```
1    0
2    0
```

```

3      0
4      0
5      0
..
57     0
58     0
59     0
60     0
61     0
Length: 61, dtype: int64

```

```
df[61].value_counts()
```

```

M      111
R       97
Name: 61, dtype: int64

```

```

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df[61] = le.fit_transform(df[61])

```

```
df.head()
```

	1	2	3	4	5	6	7	8
9 \								
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601
0.3109								
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481
0.3337								
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771
0.5598								
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276
0.0598								
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467
0.3564								
	10 ...	52	53	54	55	56	57	58
\								
0	0.2111 ...	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180	0.0084
1	0.2872 ...	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140	0.0049
2	0.6194 ...	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316	0.0164
3	0.1264 ...	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050	0.0044
4	0.4459 ...	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072	0.0048
	59	60	61					
0	0.0090	0.0032	1					

```

1  0.0052  0.0044  1
2  0.0095  0.0078  1
3  0.0040  0.0117  1
4  0.0107  0.0094  1

```

```
[5 rows x 61 columns]
```

```
df[61].value_counts()
```

```

0    111
1     97
Name: 61, dtype: int64

```

```
from sklearn.tree import DecisionTreeClassifier
```

```

y=df.iloc[:,-1]
x=df.drop([61],axis=1)

```

```

from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(x,y,test_size=0.2,random_state=25)

```

```

model=DecisionTreeClassifier()
model.fit(X_train,Y_train)
train_pred=model.predict(X_train)
test_pred=model.predict(X_test)
print(train_pred)

```

```

[0 1 0 1 1 1 0 0 1 0 1 0 0 1 1 1 1 0 1 0 1 0 0 1 1 0 0 0 1 0 0 0 1 1 0
0 0
0 0 1 0 1 0 0 1 0 0 1 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 1 1 0 1 1 1 0 0 1
1 1
0 0 1 0 1 1 0 1 1 0 1 1 1 0 0 1 1 1 0 1 0 1 0 1 1 1 0 0 1 1 0 1 1 0 1
0 0
1 0 0 1 1 1 1 1 1 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 1 1 0 0 1 0 1
1 0
0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0]

```

```

from sklearn.metrics import accuracy_score
print(accuracy_score(Y_train,train_pred))
print(accuracy_score(Y_test,test_pred))

```

```

1.0
0.7380952380952381

```

```

from sklearn.ensemble import RandomForestClassifier
clf=RandomForestClassifier(n_estimators=50)
models=clf.fit(X_train,Y_train)
train_random_pred=models.predict(X_train)
test_random_pred=models.predict(X_test)

```

```
print(accuracy_score(Y_train,train_random_pred))
print(accuracy_score(Y_test,test_random_pred))
```

```
1.0
0.8571428571428571
```

```
feature_col=list(x.columns)
feature_imp =
pd.Series(clf.feature_importances_,index=feature_col).sort_values(ascending=False)
feature_imp
```

```
11    0.071702
9     0.050912
10    0.043391
12    0.032228
47    0.030840
45    0.028840
36    0.024718
28    0.024538
13    0.024148
21    0.024134
17    0.023270
49    0.022797
2     0.021733
32    0.021601
51    0.021025
4     0.020827
16    0.020419
14    0.019385
20    0.019162
22    0.018935
48    0.018051
1     0.016490
37    0.016347
18    0.015807
5     0.015693
46    0.015113
34    0.013571
6     0.013567
31    0.013193
43    0.012970
15    0.012771
23    0.012378
39    0.012335
41    0.011997
35    0.011525
59    0.011359
60    0.011224
53    0.011162
40    0.010850
```

```

24      0.010804
55      0.010724
7       0.010521
52      0.010481
29      0.010414
19      0.010264
30      0.009630
27      0.009555
44      0.009481
50      0.009272
54      0.009113
57      0.009112
38      0.008506
58      0.008469
8       0.008399
42      0.007833
3       0.007375
33      0.006811
26      0.005852
56      0.003641
25      0.002735
dtype: float64

```

```
df.drop([25,56,26,33,3],axis=1,inplace=True)
```

```
df.drop([42,8,58,38,57],axis=1,inplace=True)
```

```
y-fi=df.iloc[:,-1]
```

```
x-fi=df.drop([61],axis=1)
```

```
X_train-fi,X_test-fi,Y_train-fi,Y_test-fi=train_test_split(x-fi,y-fi,t
est_size=0.2,random_state=30)
```

```
model=DecisionTreeClassifier()
```

```
model.fit(X_train-fi,Y_train-fi)
```

```
train_pred-fi=model.predict(X_train-fi)
```

```
test_pred-fi=model.predict(X_test-fi)
```

```
print(train_pred-fi)
```

```

[0 0 1 0 0 1 0 0 1 1 1 0 1 0 1 1 1 0 0 0 1 1 0 1 1 0 0 0 0 1 1 1 0 1 0
1 1
 1 0 1 1 1 1 1 0 0 1 1 0 0 1 1 1 0 0 0 0 1 0 1 0 1 1 0 0 1 1 0 0 0 1 1
1 0
 0 0 1 0 1 0 1 0 0 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 0 1 1 0 1 1 1 1 1 1
1 1
 0 0 0 1 0 0 1 0 0 0 1 1 1 0 1 1 0 0 1 0 1 0 1 1 0 1 0 1 1 1 0 0 0 0 0
1 0
 0 0 0 1 0 1 0 0 0 0 1 0 0 0 1 0 0 1]

```

```
from sklearn.metrics import accuracy_score
```

```
print(accuracy_score(Y_train-fi,train_pred-fi))
```

```
print(accuracy_score(Y_test-fi,test_pred-fi))
```

1.0

0.6904761904761905

```
from sklearn.ensemble import RandomForestClassifier
clf=RandomForestClassifier(n_estimators=150)
models=clf.fit(X_train_fi,Y_train_fi)
train_random_pred_fi=models.predict(X_train_fi)
test_random_pred_fi=models.predict(X_test_fi)
print(accuracy_score(Y_train_fi,train_random_pred_fi))
print(accuracy_score(Y_test_fi,test_random_pred_fi))
```

1.0

0.8571428571428571